

Options for Connecting Hadoop to Oracle

There's a comprehensive comparison matrix made by Tanel

http://www.ooug.org/wp-content/uploads/2016/05/Tanel_Poder_Connecting_Hadoop_and_Oracle.pdf

which is also available on video <https://www.youtube.com/watch?v=gSECYOi-9eA>

Below is the matrix comparison from Tanel's deck. I added the License column for completeness.

Here's the summary:

- You can load and unload data for FREE. This can be achieved by loading/reading the Hadoop data using ODBC Gateway and unload data using Sqoop
 - PRO
 - Free
 - CONS
 - ODBC gateway SQL is a serial session so if you are reading a big Hive table then expect slow performance
 - Complex SQLs could be an issue (nested subqueries, etc.) and should be pushed down to the Hive database as a new table to simplify the Oracle SQL
- Using Oracle "Big Data Connectors" (old way) and "Big Data SQL" (new way)
 - PRO
 - Allows parallel execution and complex SQLs with filter pushdown
 - CON
 - Big Data SQL – per disk drive license (just like Exadata Storage Cells)
- Using Gluent
 - PRO
 - Transparently access Oracle + Hadoop table (hybrid view) with complex SQLs. Same object name is used, that means no app change
 - Allows parallel execution on data Load/Offload with filter pushdown
 - Management tools included to maintain hybrid views and schedule offload
 - CON
 - 3rd party software – per processor license (only on Hadoop side)

Oracle <-> Hadoop data plumbing tools - capabilities

Tool	LOAD DATA	OFFLOAD DATA	ALLOW QUERY DATA	OFFLOAD QUERY	PARALLEL EXECUTION	License
Sqoop	Yes	Yes			Yes	Free
Oracle Loader for Hadoop	Yes				Yes	Part of "Big Data Connectors" Processor License
Oracle SQL Connector for HDFS	Yes		Yes		Yes	Part of "Big Data Connectors" Processor License
ODBC Gateway	Yes		Yes	Yes		Included in RDBMS license
Big Data SQL	Yes		Yes	Yes	Yes	Per Disk Drive License
Gluent Offload Engine	Yes	Yes	Yes	Yes	Yes	Processor License

Below matrix shows where the overhead is when doing read/load/unload operations:

Oracle <-> Hadoop data plumbing tools - overhead

Tool	Load data to use	Decompress CPU	Filtering CPU	Datatype Conversion	
Sqoop	Oracle	Hadoop	Oracle	Oracle	
Oracle Loader for Hadoop	Oracle	Hadoop	Oracle	Hadoop	
Oracle SQL Connector for HDFS			Oracle	Oracle / Hadoop *	Parse text files on HDFS or read pre-converted DataPump binary
ODBC Gateway		Hadoop	Hadoop	Oracle	OK for occasional archive query
Big Data SQL		Hadoop	Hadoop	Hadoop	Oracle 12c + BDA + Exadata only
Gluent Offload Engine		Hadoop	Hadoop	Hadoop	Oracle 11g,12c, no platform restrictions

Here are the other references:

- Integration of Oracle and Hadoop at CERN: hybrid databases affordable at scale
 - https://canali.web.cern.ch/canali/docs/Integration_Oracle_Hadoop_LC_CHEP2016.pdf
- Oracle Big Data Connectors
 - <https://www.abis.be/resources/presentations/obug20130326bigdata.pdf>
- Oracle and Hadoop, let them talk together
 - https://www.doag.org/formes/pubfiles/9479832/2017-DA-Laurent_Leturgez-Hadoop_and_Oracle_Let_Them_Talk_Together_-Manuskript.pdf
- How To Query Hive and Impala from Oracle using ODBC Heterogeneous Gateway
 - <https://community.oracle.com/docs/DOC-1002634>
- Oracle Big Data SQL - copy2hadoop
 - <https://blogs.oracle.com/datawarehousing/data-loading-into-hdfs-part2-data-movement-from-the-oracle-database-to-the-hdfs>
 - <https://docs.oracle.com/bigdata/bds31/BDSUG/copy2bda.htm#BIGUG76738>

Price List:

- <https://www.oracle.com/assets/technology-price-list-070617.pdf>

	Named User Plus	Software Update License & Support	Processor License	Software Update License & Support
Berkeley DB - Transactional Data Store	-	-	5,800	1,276.00
Berkeley DB - Concurrent Data Store	-	-	1,800	396.00
Berkeley DB - Data Store	-	-	900	198.00
Berkeley DB Java Edition - High Availability	-	-	9,800	2,156.00
Berkeley DB Java Edition - Transactional Data Store	-	-	5,800	1,276.00
Berkeley DB Java Edition - Concurrent Data Store	-	-	1,800	396.00
Berkeley DB XML - High Availability	-	-	13,800	3,036.00
Berkeley DB XML - Transactional Data Store	-	-	8,100	1,782.00
Berkeley DB XML - Concurrent Data Store	-	-	2,600	572.00
Berkeley DB XML - Data Store	-	-	1,800	396.00

Other Products

Audit Vault and Database Firewall	120	26.40	6,000	1,320.00
Big Data Connectors	-	-	2,000	440.00
Big Data Spatial and Graph	-	-	2,000	440.00

	License Price	Software Update License & Support	Licensing Metric
Key Vault	100,000	22,000.00	Per Server
Big Data SQL	4,000	880.00	Per Disk Drive

- Big Data Connectors User's Guide
 - https://docs.oracle.com/cd/E27101_01/doc.10/e27365/start.htm
- Big Data SQL Licensing
 - <https://docs.oracle.com/bigdata/bds31/BDSIG/licensing.htm#BDSIG-GUID-8E911C46-3FAA-44BA-BCFE-FB9EF94E494E>
- Database Gateway and Generic Connectivity (DG4ODBC) Licensing Considerations (Doc ID 232482.1)

